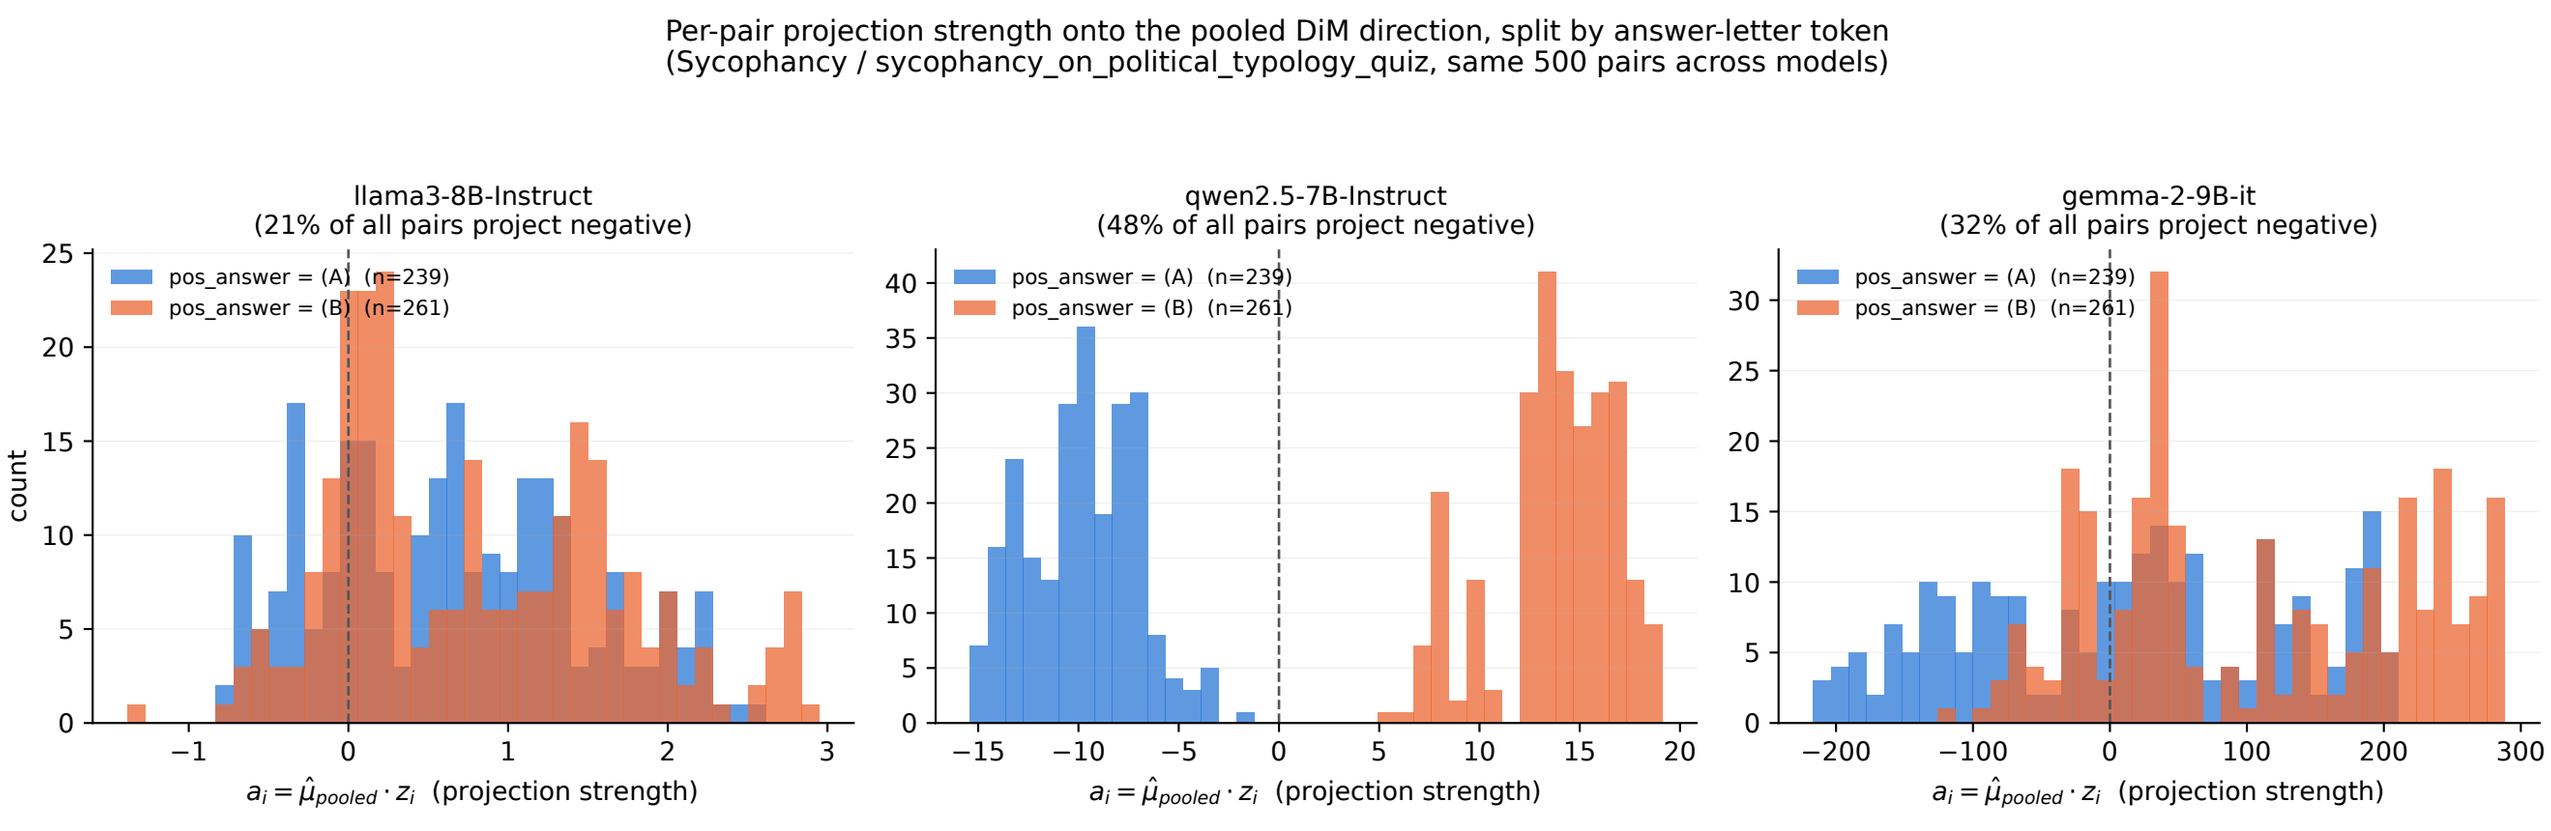
Per-pair projection strength onto the pooled DiM direction, split by answer-letter token
(Sycophancy / sycophancy_on_political_typology_quiz, same 500 pairs across models)
llama3-8B-Instruct
(21% of all pairs project negative)
qwen2.5-7B-Instruct
(48% of all pairs project negative)
gemma-2-9B-it
(32% of all pairs project negative)
pos_answer = (A) (n=239)
pos_answer = (B) (n=261)
count
$a_i = \hat{\mu}_{pooled} \cdot z_i$ (projection strength)