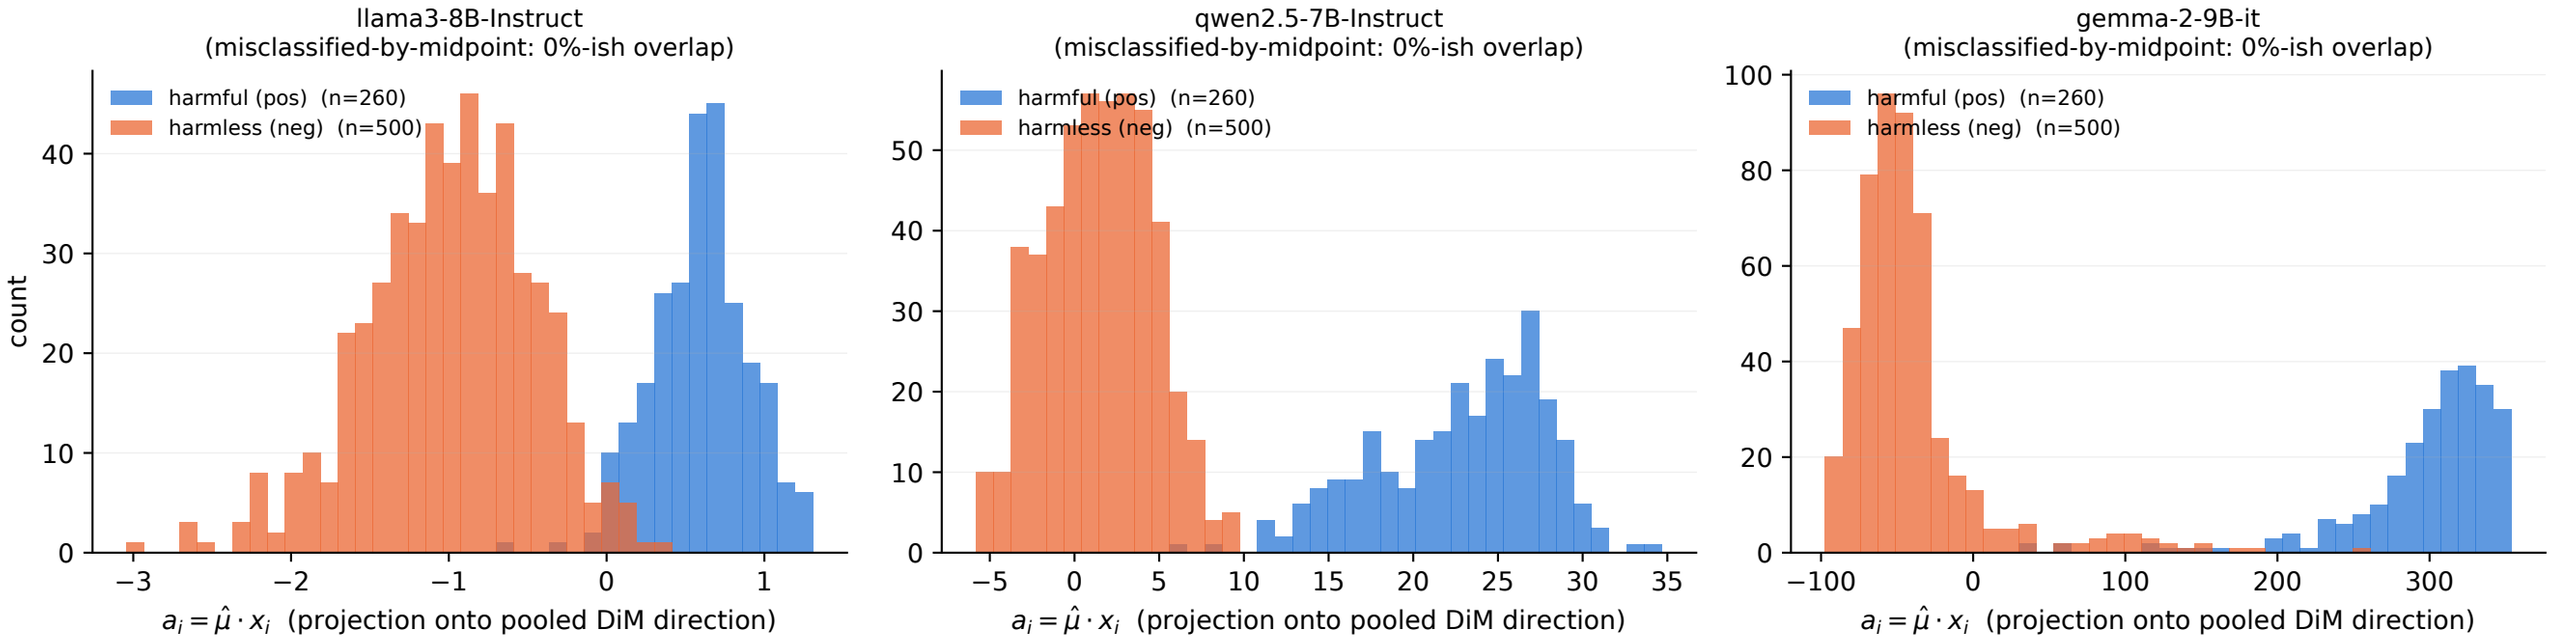
Per-sample projection strength of harmful (pos) vs. harmless (neg) activations
onto their own pooled DiM direction (Refusal dataset)
llama3-8B-Instruct
(misclassified-by-midpoint: 0%-ish overlap)
qwen2.5-7B-Instruct
(misclassified-by-midpoint: 0%-ish overlap)
gemma-2-9B-it
(misclassified-by-midpoint: 0%-ish overlap)
harmful (pos) (n=260)
harmless (neg) (n=500)
count
$a_i = \hat{\mu} \cdot x_i$ (projection onto pooled DiM direction)